

WEB SCRAPING FOR AI TRAINING: HOW BIG TECH COMPANIES BENEFIT FROM PUBLIC DATA

Vinícius Queiroz Soares

Keywords: AI, Web Scraping, Open Source, LLM, Pre-training, Data Pipeline

Abstract

The rapid advancement of Artificial Intelligence is predicated on the consumption of massive, publicly available web archives. This paper provides a comprehensive analysis of the data processing pipelines that transform raw Common Crawl data into training-ready corpora, exemplified by C4, FineWeb, and Nemotron-CC. Through a technical review of extraction, deduplication, and filtering methodologies, we demonstrate that the transition from heuristic "blocklists" to AI-driven synthesis multiplies existing risks, including political bias, minorities discriminations and the potential for "Model Collapse" arising from an over-reliance on synthetic data.

Beyond technical architecture, the study critically examines the legal and ethical crises defining the sector. We analyze the precarious application of the "fair use" doctrine to commercial training and evaluate Collective Rights Management (CRM) organizations as a pragmatic, market-based solution to the copyright infringement deadlock. Furthermore, we investigate the emerging "data winter" precipitated by high-quality publishers blocking AI crawlers via *robots.txt*. We conclude that this "opt-out" trend threatens to skew future datasets toward polarized, low-quality content, suggesting that the current paradigm of indiscriminate scraping is unsustainable and necessitates new frameworks for data governance and creator compensation.

1. Introduction

Recent years have witnessed remarkable progress in Artificial Intelligence (AI), largely driven by pre-trained models whose capabilities are strongly correlated with the size of

their pre-training corpora (Wenzek et al., 2019; Raffel et al., 2020). Assembling datasets on the required scale of trillions of tokens is a significant challenge, as high-quality curated sources like Wikipedia are limited in size and linguistic diversity (Dodge et al., 2021). Consequently, the vast, unstructured data available on the public web has become an essential resource, positioning the non-profit Common Crawl archive as critical infrastructure for the entire generative AI ecosystem (Baack, 2024).

However, the raw data from Common Crawl is not a ready-to-use dataset. It must undergo extensive processing via complex pipelines, such as those that produce C4 (Raffel et al., 2020), CCNet (Wenzek et al., 2019), and Nemotron-CC (Su et al., 2025). These pipelines are designed to deduplicate, filter, and clean the data, but the specific methodologies have profound, often underexamined, ethical implications. For example, filtering techniques intended to remove offensive language have been shown to disproportionately exclude content from marginalized communities (Dodge et al., 2021).

This reliance on web-scraped data, even when mediated by non-profits, places commercial AI development in a legally and ethically precarious position. Tech companies benefit from this data while facing increasing legal challenges over copyright infringement, with the "fair use" doctrine being heavily contested in court (Kim, 2024; Opderbeck, 2024). Simultaneously, a new challenge has emerged: high-quality data sources, particularly news organizations, are actively blocking AI crawlers using robots.txt (Fletcher, 2024). This "opt-out" trend threatens to skew future training datasets toward lower-quality, hyper-partisan content, creating a significant risk of embedding these biases into foundational models (Bouchaud & Ramaciotti, 2025).

This paper provides a comprehensive analysis of the end-to-end pipeline from raw web crawl to usable AI training data. It synthesizes the technical processes, investigates the infrastructural role of public data, and critically examines the intersecting legal and ethical crises that define this domain. We analyze the specific data processing pipelines, the legal arguments surrounding fair use, and the emerging threat of data

skew from crawler restrictions.

2. Methodology

This study employed a mixed-methods approach to analyze the role of web-scraped data in AI training. The research was conducted from late 2025 to early 2026 and combined a technical analysis of data-processing pipelines with a systematic review of legal and ethical frameworks.


The first stage involved a technical review of publicly documented web data archives and their derivatives. We analyzed the structure and limitations of the Common Crawl archive (Baack, 2024) and synthesized the processing stages of prominent derivative datasets, including C4 (Raffel et al., 2020), CCNet (Wenzek et al., 2019), and Nemotron-CC (Su et al., 2025). This included an examination of scraping techniques (e.g., static vs. dynamic extraction) and the specific filtering, deduplication, and quality-scoring methods employed.

The second stage comprised a systematic review of legal and ethical scholarship and contemporary case law. We focused on the application of the fair use doctrine to AI training, particularly in light of recent jurisprudence (Kim, 2024; Opderbeck, 2024). This was supplemented by an analysis of ethical issues, including data privacy (Kahlon & Singh, 2024), algorithmic bias, and the documented harms of data filtering (Dodge et al., 2021).

Finally, we analyzed the emerging trend of AI crawler restrictions. We synthesized recent findings on the adoption of robots.txt rules to block crawlers (Fletcher, 2024) and the potential for this trend to introduce systemic bias into future AI datasets (Bouchaud & Ramaciotti, 2025).

The study's analysis of web data processing stages is complemented by a practical implementation found in a public repository. Key extraction and filtering techniques were demonstrated and are available in a public repository, which includes annotated scripts

and sample data snippets illustrating the processing stages discussed in this paper: [access to repository](#).



The web crawling project consists of a two-stage data processing pipeline designed for generating high-quality AI training datasets. The first stage involves the Puppeteer-based WikipediaCrawler, which handles dynamic content extraction and supports both single-page and recursive crawling with configurable limits. This raw data is immediately filtered using jusText integration, serving as the primary mechanism to remove boilerplate content (navigation, headers, footers) and ensure only clean, main article text is retained.

The second stage introduces the C4-like DataPipeline, which implements advanced quality-scoring by applying a range of heuristics to the extracted text. These heuristics check for quality based on content length, structure (e.g., Punctuation ratios), and

composition (e.g., limits on Symbols & Digits or Uppercase letters). Additionally, the pipeline performs Language Identification for filtering by target languages (e.g., 'eng') and incorporates robust N-gram Deduplication (defaulting to 13-word n-grams and an 80% overlap threshold) to prevent the inclusion of duplicate or near-duplicate content, thereby producing clean, high-variety data suitable for training large language models.

3. Findings and Analysis

Our analysis reveals a deeply symbiotic, yet tension-filled, relationship between public web data and large-scale AI development. The findings are presented in three parts: the technical data pipeline, the legal and ethical conflicts, and the emerging challenge of data access.

3.1 The Data-to-Corpus Pipeline: From Common Crawl to C4 and beyond

The foundational layer for modern Large Language Models (LLMs) is Common Crawl, a massive, open repository of web data typically accessed via WET files (text-only extractions). While this corpus serves as the training bedrock for state-of-the-art models, it is not a neutral or representative mirror of the internet.

As noted by Baack (2024), the dataset is constrained by significant structural limitations. The crawling methodology relies on **harmonic centrality** (link density) to prioritize URLs, which inadvertently overrepresents mainstream domains while marginalizing niche communities and minority languages.

Furthermore, the raw data is compromised by quality issues, including high volumes of SEO spam, duplicates, and "bad" HTML. Consequently, raw Common Crawl data requires an extensive, multi-stage "funnel" of processing to be viable for model training.

3.2 Evolution of Processing Methodologies

The transformation of raw web scrapes into high-fidelity training tokens has evolved through three distinct generations of engineering, moving from rigid heuristics to

AI-driven synthesis.

3.2.1 First Generation: Heuristic Filtering (C4)

The Colossal Clean Crawled Corpus (C4) established the initial standard for data hygiene by applying a strict, heuristic-based pipeline to a single Common Crawl snapshot. This methodology prioritized precision through rigid rules:

- **Language and Punctuation Constraints:** The pipeline utilized `langdetect` to enforce a 99% English probability threshold. A controversial "punctuation rule" was applied, discarding any line not ending in terminal punctuation (., ?, !). While effective for prose, this heuristically removed significant amounts of code and creative text.
- **Exclusionary Blocklisting:** C4 filtered pages containing words from a "bad words" blocklist. This approach acts as a blunt instrument; it often removes safe, identity-related content from marginalized groups while failing to catch toxic content that avoids specific banned terminology.

3.2.2 Second Generation: Advanced Extraction and Fuzzy Deduplication (FineWeb)

FineWeb (2024) represents a paradigm shift toward higher data retention and sophisticated engineering, processing 96 snapshots to maximize volume. The pipeline replaced basic WET file parsing with the *trafilatura* library, significantly improving the extraction of main content by stripping HTML boilerplate such as menus and footers.

A critical advancement in FineWeb is the implementation of MinHash-based fuzzy deduplication. Unlike C4's line-level exact matching, MinHash identifies documents that are semantically identical despite minor variations (e.g., templated news or syndicated articles).

This generation also introduced the FineWeb-Edu subset, which utilizes an LLM

(Llama-3) to grade content quality, marking a departure from keyword-only filtering. However, this rigorous filtering results in the removal of approximately 90% of the raw data, creating a trade-off between data quality and the token volume required for long-horizon training.

3.2.3 Third Generation: AI-Driven Curation and Synthesis (Nemotron-CC)

The current state-of-the-art, Nemotron-CC (2025), replaces static heuristics with dynamic, AI-native judgments. This approach addresses the limitations of rule-based filtering through two key mechanisms:

- **Classifier Ensembling:** Rather than relying on binary blocklists, Nemotron-CC employs a voting system comprised of multiple AI classifiers to assess document quality. This nuanced approach preserves safe but non-standard text that traditional heuristics often discard.
- **Synthetic Data and Hybrid Deduplication:** The dataset is enriched with 1.9 trillion synthetic tokens. These are AI-generated to rephrase or "repair" lower-quality data, effectively synthesizing new, high-quality training signals. Additionally, the pipeline utilizes a hybrid deduplication strategy, combining global fuzzy deduplication with exact substring matching to ensure maximum data novelty.

Comparative Summary

Methodology	C4 (2019)	FineWeb (2024)	Nemotron-CC (2025)
Filtering Logic	Strict Heuristics (Rules)	Fuzzy Dedup & Classifiers	AI Ensembling & Synthesis

Deduplication	Line-level (Exact)	Document-level (MinHash)	Global Fuzzy + Substring
Quality Control	"Bad Word" Lists	Educational Classifiers (Llama-3)	Model Ensembling + Synthetic Repair
Data Composition	1 Snapshot	96 Snapshots	Hybrid (Crawl + Synthetic)

3.3 Why Deduplication Matters?

Deduplication serves as a fundamental preprocessing filter in the construction of large-scale datasets for Large Language Models (LLMs). By removing redundant data, this process ensures that models learn to generalize language patterns rather than memorizing and regurgitating specific, over-represented text strings.

3.3.1 The Necessity and Impact of Data Cleaning

The presence of duplicate content creates significant risks for model training, including the failure to generalize and the potential for copyright or privacy violations through verbatim data reproduction. In the specific context of the CCNet pipeline, Wenzek et al. (2019) demonstrate that duplicated paragraphs are pervasive, representing approximately 70% of the text in Common Crawl snapshots.

Implementing deduplication provides two distinct advantages for dataset quality:

- **Removal of Noise:** The process eliminates "boilerplate" content common in web

crawling, such as navigation menus, cookie warnings, and contact information.

- Improved Language Identification: Since much of this boilerplate is English text found on pages written in other languages, its removal significantly improves the robustness of Language Identification (LID). This is particularly critical for maintaining the integrity of datasets for low-resource languages⁴.

3.3.2 Methodologies for Deduplication

Deduplication techniques generally rely on hashing algorithms to identify identical or similar content efficiently, avoiding the computational prohibitive cost of direct text comparison.

3.3.2.1 Paragraph-Level Hashing (Exact Match)

The methodology employed by Wenzek et al. (2019) focuses on identifying exact paragraph matches. The pipeline first normalizes text by lower-casing characters, replacing numbers with placeholders, and removing Unicode punctuation and accents⁵. Following normalization, the system computes a unique key for each paragraph using the first 64 bits of the SHA-1 digits⁶. These keys are compared against binary files of previously processed hashes to identify and discard duplicates.

3.3.2.2 Near-Duplicate Detection (MinHash and LSH)

To identify documents that are "mostly the same" but not identical (near-duplicates), researchers often utilize a pipeline designed to estimate similarity probabilistically. This approach creates a compact "fingerprint" for documents to approximate similarity metrics without comparing full text.

This process generally operates in three stages:

- Shingling: Documents are decomposed into sets of overlapping n-grams (shingles). The similarity between two documents is defined by the Jaccard

Similarity:

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

- MinHash: To reduce the dimensionality of the data, the MinHash algorithm creates a fixed-size signature for each document. By selecting the minimum hash value from multiple hash functions, the system generates a signature that preserves the Jaccard similarity estimate while drastically reducing storage requirements.
- Locality-Sensitive Hashing (LSH): To solve the computational bottleneck of comparing every signature to every other signature $O(n^2)$, LSH is employed. This technique places similar signatures into the same "buckets," ensuring that only documents colliding within a bucket are compared. This reduces the complexity to roughly $O(n)$, making the deduplication of trillions of tokens computationally feasible.

3.4 Language Identification and Quality Filtering

Data processing in natural language processing (NLP) pipelines fundamentally serves as a gatekeeping mechanism, ensuring that training corpora are both linguistically appropriate and of sufficient quality. This stage is bipartite, consisting of language identification followed by rigorous quality filtering.

3.4.1 Vector-Based Language Identification

The foundational premise of language identification is the recognition of unique "statistical fingerprints" inherent to specific languages. While classical approaches relied on counting character N-grams (e.g., unigrams, bigrams), modern pipelines have largely transitioned to vector-based models, such as *fastText*.

Unlike simple N-gram frequency counters, vector-based models represent character N-grams as continuous vectors embedded in a high-dimensional space. In this schema, linguistically similar texts cluster together. The model identifies the language by determining the dominant vector direction of the input text. This approach offers significant advantages in robustness and speed; fastText is capable of classifying thousands of documents per second and maintains high accuracy even when processing noisy data or texts exhibiting code-switching (e.g., "I'm going to the *tienda*").

3.4.2 Approaches to Quality Filtering

Once language is established, the pipeline must excise low-quality data—defined as boilerplate, spam, or incoherent text—to prevent the degradation of model performance. Two primary methodologies have emerged to address this challenge: perplexity-based filtering and model-based classification.

3.4.2.1 Perplexity-Based Filtering (Raffel et al., 2020)

This heuristic approach, widely utilized in the construction of datasets such as C4, predicates quality on the statistical concept of "surprise." It posits that high-quality text follows a predictable probability distribution similar to a reference corpus, typically Wikipedia. To quantify this, we calculate the perplexity of a token sequence as the inverse probability of the text, normalized by the sequence length.

In practice, a lightweight model (such as KenLM) computes this score. Texts with low perplexity are deemed "unsurprising" and retained, while those with high perplexity—indicating disjointed or garbled content—are discarded. However, this method functions as a blunt instrument; it frequently misclassifies creative writing, poetry, or dialectal variations as "noise" simply because they statistically deviate from the encyclopedic tone of the reference corpus.

3.4.2.2 Model-Based Classifiers (Su et al., 2025)

Recent advancements, exemplified by datasets like FineWeb and Nemotron-CC, employ a more sophisticated, supervised learning approach. Rather than relying on statistical heuristics, researchers train binary classifiers (often Transformers) to distinguish between "high-quality" and "low-quality" content directly.

The training process involves curating a positive dataset of "gold standard" texts (e.g., textbooks, peer-reviewed articles) and a negative dataset of random web noise. The resulting model outputs a scalar probability score (e.g., 0.98), allowing for nuanced, targeted filtering. This method enables curators to define quality specifically—for instance, prioritizing educational value—rather than relying on a generic definition of grammatical correctness.

3.4.3 Bias and Discriminatory Impact

While filtering is critical for computational efficiency and model coherence, it introduces significant ethical risks regarding representation. Dodge et al. (2021) conducted a critical analysis of "blocklist" filtering—a standard preprocessing step where documents containing specific terms are automatically removed.

Their findings demonstrated that these blocklists are not semantically neutral but exert a discriminatory impact on marginalized communities. Specifically, the study revealed that content written in African American English (AAE) and texts discussing LGBTQ+ identities were disproportionately filtered out. Because this removal occurs prior to training, the resulting models exhibit a systematic bias toward "White Mainstream English" and lack the contextual knowledge necessary to generate accurate or safe text relevant to these excluded groups. This underscores the tension between data sanitation and the preservation of linguistic diversity.


3.5 Synthetic Augmentation

The evolution of training methodologies has necessitated a shift from passive data

collection to active data refinement, a process exemplified by recent frameworks such as Nemotron-CC (Su et al., 2025). This approach addresses the inherent deficiencies of raw web data—which, while abundant, is frequently comprised of structurally poor, typo-laden, or low-fidelity content. In this paradigm, the limitations of original scraped data are not viewed as a dead end, but rather as a substrate for improvement through synthetic augmentation.

As demonstrated in the Nemotron-CC methodology, powerful "teacher" models are utilized to process this noisy input. By instructing the model to rephrase, reorganize, or rewrite the original messy text into high-quality articles, developers can synthetically elevate the utility of the dataset. This process ensures that the core informational value is retained while the noise and irregularities common in raw web scrapes are stripped away. Consequently, the training corpus is transformed from a collection of disparate, lower-quality samples into a coherent, high-fidelity dataset that is significantly more valuable for instruction-tuning and model robustness than the original source material.

While synthetic data offers a solution to scarcity, an over-reliance on such generated content introduces a critical failure mode known as "Model Collapse," occasionally referred to in literature as the "Habsburg Jaw" effect. This phenomenon describes a degenerative process that occurs when models are iteratively trained on the outputs of previous model generations rather than on ground-truth human data.



Much like the progressive degradation of detail observed when creating a photocopy of a photocopy, recursive training on synthetic data causes the model to lose the fidelity of the original distribution. As the model "forgets" the inherent diversity, nuance, and irregularity—or "weirdness"—of authentic human language, its outputs converge toward a bland statistical average. This convergence results in a loss of variance, where the model becomes repetitive, less creative, and increasingly prone to errors. Consequently, to maintain model robustness and prevent this degeneration, synthetic data must be utilized strictly as a supplement to high-quality human data, rather than a complete replacement.

3.6 Legal Analysis: Fair Use and Commercial AI

The central legal conflict regarding the training of Artificial Intelligence (AI) models is rooted in the statutory definition of copyright infringement. As Kim (2024) articulates, the technical architecture of Large Language Models (LLMs) necessitates a process that inherently conflicts with the exclusive right of reproduction. This infringement is not merely an incidental outcome but a fundamental step in the data ingestion pipeline; web scrapers must retrieve, copy, and store copyrighted text and images to generate the

requisite training datasets. While the resulting AI models are comprised of uncopyrightable mathematical tokens, the initial digitization and retention of the original expressive works constitute a prima facie violation of copyright law. The persistence of these copies in training databases undermines defenses regarding the temporary nature of the reproduction, suggesting that the infringement is both direct and systematic.

To counter these claims of infringement, AI developers heavily rely on the affirmative defense of fair use. This legal doctrine involves a balancing test designed to accommodate new technologies without stifling the rights of creators. The analysis is governed by four statutory factors: the purpose and character of the use, the nature of the copyrighted work, the amount and substantiality of the portion used, and the effect of the use upon the potential market for the original work. In the context of AI, the adjudication of these factors determines whether the massive scale of copying serves a distinct, legally protected purpose or merely exploits the labor of rights holders without compensation.

The application of the first factor—the purpose and character of the use—often turns on the concept of "transformativeness." Opderbeck (2024) explains that AI companies argue their ingestion of data is transformative because it fundamentally alters the purpose of the original work. Under this theory, the AI does not reproduce the work to convey its original expressive intent, but rather analyzes it to extract statistical correlations and functional patterns. By converting creative expression into computational data, proponents argue that the training process creates a new utility that is distinct from the original consumption of the media. However, this argument is increasingly challenged by the reality that the "input" retains the essential qualities of the copyrighted expression, and the "output" may compete directly with the original works in the marketplace.

3.7 Copyright Infringement: CRMs as a potential solution

The primary challenge in licensing data for AI training is the "prohibitive transaction costs" associated with the massive scale of data involved; it is impossible to negotiate with billions of individual rights holders. The text proposes Market Clearinghouses and Collective Rights Management (CRM) organizations as the most pragmatic, market-based solution. These entities serve as a form of "private ordering" to manage this complexity. Their core functions are to:

- Aggregate copyrights from vast numbers of creators.
- Offer standardized licenses to AI developers.
- Manage the collection and distribution of licensing revenues back to the rights holders.

This strategy is based on successful existing models in other creative industries. There are specific examples of organizations that already simplify complex licensing, demonstrating that this framework is viable. These models include:

- Traditional CRMs (e.g., ASCAP, BMI): These Performance Rights Societies manage music catalogs, allowing venues to purchase a single "blanket license" to play millions of songs.
- Patent Pools: These organizations bundle patents for core technologies to streamline development and manufacturing.
- Market Clearinghouses (e.g., Getty Images): Getty is a key example, aggregating and licensing visual media. The text notes it is a "natural extension" for Getty to license its tagged catalog for AI training, a right the company is already asserting through legal action.

These models could be readily adapted for the two main types of AI training data. For commercially available content (books, music, films), CRMs could coordinate with existing distributors like Amazon or major record labels. For individual content from social media, platforms like YouTube or TikTok could update their terms of service to

include AI training and develop a revenue-sharing model with their users. The primary benefits of this approach are providing compensation to creators and improving data integrity, as market forces would incentivize clean, high-quality datasets.

Despite being a practical solution, there are significant trade-offs and obstacles. While they solve transaction costs, these organizations add their own administrative fees and raise new concerns. The key challenges include:

- **Anticompetitive Behavior:** Large aggregators, like traditional CRMs, can face antitrust scrutiny for using their market power to "squeeze both content contributors and consumers."
- **Compensation Difficulty:** While creators would be paid, the sheer volume of data means that payments to "individual creators" might not be significant.
- **Logistical & Legal Hurdles:** The scale of web data far exceeds any existing CRM's repertoire. In the EU, specific "opt-out" rules and the high risk of fraudulent actors impersonating rights holders create massive barriers that would be difficult to overcome without government backing.

3.8 Ethical Analysis: Bias, Privacy, and Data Laundering

The technical necessity of data filtering inevitably intersects with significant ethical risks regarding representation and fairness. As demonstrated by Dodge et al. (2021), the standard "blocklist" filtering techniques employed in datasets like C4 are not semantically neutral but actively discriminatory. By automatically excising documents containing specific terms deemed "offensive," these pipelines disproportionately remove non-toxic content generated by and relevant to marginalized communities. Specifically, the study revealed that content written in African American English (AAE) and texts discussing LGBTQ+ identities were filtered out at significantly higher rates than "White Mainstream English," leading to allocational harms where resulting models lack the contextual knowledge necessary to serve these groups effectively.

Beyond the immediate issue of bias, the scraping of public data raises profound privacy

concerns. Kahlon & Singh (2024) and Baack (2024) highlight that the indiscriminate harvesting of the web captures vast amounts of Personally Identifiable Information (PII) without consent. Unlike curated datasets where consent is explicit, web archives inadvertently store private data that becomes embedded in model weights, making "the right to be forgotten" technically difficult to enforce once training is complete.

Furthermore, the ecosystem relies on a structural dynamic described by Kim (2024) as "AI data laundering." In this framework, non-profit entities like Common Crawl absorb the initial legal and reputational risks of massive data collection under the guise of archival research. Commercial tech companies then utilize these "laundered" datasets for profit, effectively externalizing the liability of copyright infringement and privacy violations while internalizing the commercial benefits. This separation of collection and utilization complicates accountability mechanisms, as commercial developers can claim they are merely using public, open-source archives rather than directly scraping protected content.

3.9 The robots.txt Dilemma and Dataset Skew

A rapidly accelerating challenge to the viability of future datasets is the widespread adoption of crawler restrictions by premium content providers. Fletcher (2024) documents a significant trend where high-quality news organizations and academic repositories are increasingly updating their robots.txt files to explicitly block AI scrapers. This "opt-out" movement is driven largely by the lack of compensation for intellectual property and the competitive threat posed by generative search engines, effectively walling off the most factual, fact-checked, and linguistically rigorous portions of the public web from future training runs.

The implications of this exclusion go beyond a mere reduction in data volume; they introduce a systemic "dataset skew." As analyzed by Bouchaud & Ramaciotti (2025), as high-quality, neutral sources actively withdraw from the training pool, the remaining accessible data is increasingly comprised of lower-quality, unmoderated, and

hyper-partisan content that has no incentive to restrict access. Consequently, future Large Language Models face the risk of being trained on a polarized version of the internet, potentially amplifying existing societal biases and degrading the factual reliability of foundational models.

4. Discussion

The synthesis of technical, legal, and ethical findings indicates that the "data-cleaning" pipelines described in Section 3.1 are not merely neutral sanitation processes but active forms of curation that embed specific values and biases. The findings of Dodge et al. (2021) confirm that the pursuit of "quality" through heuristic filtering frequently comes at the cost of diversity, silencing minority voices under the guise of safety. This creates a fundamental tension in AI development: the mechanisms designed to make models safe for commercial deployment simultaneously render them less representative of the diverse user base they are intended to serve.

Furthermore, the industry is facing a potential "data winter" precipitated by the mass opt-out of high-quality publishers. If the trend identified by Bouchaud & Ramaciotti (2025) continues, tech companies will be forced to navigate a difficult trade-off: rely on a shrinking pool of high-quality human data, utilize lower-quality accessible data which aggravates bias, or lean more heavily on synthetic data as proposed by Su et al. (2025). However, as noted in the analysis of Nemotron-CC, an over-reliance on synthetic data risks "model collapse," suggesting that synthetic augmentation is a supplement, not a complete substitute, for the authentic human expression that is currently being fenced off.

Finally, a legal and ethical paradox defines the current landscape. Tech companies are caught between conflicting pressures: they face aggressive litigation for scraping copyrighted material without permission (Kim, 2024), yet when high-quality sources voluntarily opt out, the resulting "dataset skew" threatens the safety and neutrality of their products. This situation highlights the precarious position of Common Crawl

(Baack, 2024), which serves as a critical "data commons." This infrastructure is currently under threat from both legal challenges regarding copyright and the degradation of the commons itself, framing the issue not just as a technical hurdle, but as a crisis of digital governance and infrastructure.

5. Limitations

This analysis is primarily constrained by its reliance on publicly documented datasets and pipelines such as C4, CCNet, and Nemotron-CC. The proprietary data processing pipelines employed by major commercial entities like OpenAI or Google are likely more sophisticated and may differ significantly in their handling of filtering and deduplication; however, because these methodologies remain trade secrets, they cannot be fully assessed.

Additionally, the legal analysis presented is based on case law and scholarship current as of late 2025. Given the rapid pace of litigation regarding copyright and fair use, the resolution of ongoing high-profile cases could fundamentally alter the legal frameworks discussed. Finally, the analysis of "dataset skew" resulting from robots.txt restrictions relies on early findings by Bouchaud & Ramaciotti (2025); the long-term impact of these restrictions on model performance remains a projection that requires longitudinal study to fully quantify.

6. Conclusion

This paper has analyzed the end-to-end lifecycle of AI training data, tracing its journey from the raw, unstructured public web to its refinement in foundational models. The investigation reveals that the "data cleaning" pipelines described—often viewed as neutral engineering tasks—are in fact active forms of curation that embed profound legal risks and ethical biases.

Technically, we established that while the industry is transitioning from heuristic filtering (C4) to AI-driven synthesis (Nemotron-CC), this evolution introduces new vulnerabilities.

As demonstrated by Su et al. (2025), the increasing reliance on synthetic data to compensate for quality deficits risks inducing "Model Collapse", where recursive training on generated output degrades model fidelity. Furthermore, our review of filtering impacts confirms that the pursuit of "safe" data frequently results in allocational harms. As noted by Dodge et al. (2021), standard blocklists disproportionately silence marginalized communities, thereby enforcing a "White Mainstream" linguistic norm on foundational models.

Legally and infrastructurally, the ecosystem is at a tipping point. The reliance on the "fair use" doctrine is increasingly precarious (Kim, 2024), and the absence of a standardized compensation framework has triggered a "data winter." The rapid adoption of *robots.txt* restrictions by high-quality publishers (Fletcher, 2024) threatens to skew future datasets toward hyper-partisan and low-quality content (Bouchaud & Ramaciotti, 2025). While market-based solutions such as Collective Rights Management (CRM) organizations offer a viable path for licensing and compensation, their implementation faces significant logistical and antitrust hurdles.

Future Work Future research must move beyond diagnosing these issues to engineering solutions. Specifically, technical research should focus on developing "bias-aware" filtering pipelines that can distinguish between toxic content and non-standard dialects like AAE, ensuring safety does not come at the cost of inclusion. Additionally, longitudinal studies are urgently needed to quantify the long-term impact of the "opt-out" trend on model hallucination rates and political neutrality. Finally, legal scholars must further explore the feasibility of an international "Data Commons" framework that could supersede the fragmented CRM model.

Ultimately, the generative AI industry has built its massive valuation on the assumption of a free, open, and accessible web. As that assumption falters under the weight of legal challenges and crawler restrictions, this research underscores that the next generation of AI development cannot rely solely on better algorithms. It requires a fundamental restructuring of data governance, transparency, and creator compensation to preserve

the integrity of the information ecosystem upon which these models depend.

References

Baack, S. (2024). *Training Data for the Price of a Sandwich: Common Crawl's Impact on Generative AI*. Mozilla Insights.

Baack, S. (2024). A Critical Analysis of the Largest Source for Generative AI Training Data: Common Crawl. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency* (FAccT '24). ACM, New York, NY, USA, 2199–2208. <https://doi.org/10.1145/3630106.3659033>

Bouchaud, P., & Ramaciotti, P. (2025). *Web Crawler Restrictions, AI Training Datasets & Political Biases*. arXiv preprint arXiv:2510.09031v1.

Dodge, J., Sap, M., Marasović, A., Agnew, W., Ilharco, G., Groeneveld, D., Mitchell, M., & Gardner, M. (2021). *Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus*. arXiv preprint arXiv:2104.08758v2.

Fletcher, R. (2024, February 22). *How many news websites block AI crawlers?* Reuters Institute for the Study of Journalism.

Kahlon, N. K., & Singh, W. (2024). *A Systematic Review of Web Scraping: Techniques, LLM-Enhanced Approaches, Performance Metrics, and Legal-Ethical Issues*. Preprint available at SSRN: <https://ssrn.com/abstract=5429131> (Accessed October 28, 2025).

Kim, J. Y. C. (2024). Data Scraping for Generative AI—To What Extent? *Brooklyn Journal of Corporate, Financial & Commercial Law*, 19(1), Article 8.

Opderbeck, D. W. (2024). Copyright in AI Training Data: A Human-Centered Approach. *Oklahoma Law Review*, 76(4), 951-1023.

Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2020). Exploring the limits of transfer learning with a unified text-to-text

transformer. *Journal of Machine Learning Research*, 21(140), 1-67.

Su, D., Kong, K., Lin, Y., Jennings, J., Norick, B., Kliegl, M., Patwary, M., Shoeybi, M., & Catanzaro, B. (2025). *Nemotron-CC: Transforming Common Crawl into a Refined Long-Horizon Pretraining Dataset*. arXiv preprint arXiv:2412.02595v2.

Wenzek, G., Lachaux, M.-A., Conneau, A., Chaudhary, V., Guzmán, F., Joulin, A., & Grave, E. (2019). *CCNet: Extracting High Quality Monolingual Datasets from Web Crawl Data*. arXiv preprint arXiv:1911.00359v2.